

# Week 1: Potential Outcomes

PLSC 30600 - Causal Inference

Welcome!

# Course Overview

- Instructor: Anton Strezhnev
- TA: Cindy Wang
- Logistics:
  - Lectures T/Th - 2:00pm - 3:20pm; Sections F - 2:30pm - 3:20pm
  - 3 Problem Sets (~ 2 weeks)
  - Take-home midterm (~ 1 week)
  - In-class final (exact day/time TBA)
  - My office hours: Tuesdays 4pm-6pm (Pick 328)
  - Cindy's office hours: TBA
- What is this course about?
  - *Defining* causal effects w/ a structured statistical framework (potential outcomes)
  - *Outlining* assumptions necessary to identify causal effects from data
  - *Estimating* and conducting *inference* on causal effects
- Goals for the course
  - Give you the tools you need to develop your own causal research designs and comment on others' designs.
  - Equip you with an understanding of the fundamentals of causal inference to enable you to learn new methods.
  - Professionalization -- how did causal inference emerge as a field and how do different disciplines approach the topic.

# Course workflow

- Lectures
  - Topics organized by week
  - Tuesday lecture: Introduction + higher-level overview
  - Thursday lecture: More in-depth details + applications
  - You should do the readings prior to each week, but definitely be sure to do them before the Thursday.
- Readings
  - Mix of textbooks and papers - theory + application.
  - Organized on the syllabus by priority. Try to do all of them, but if you need to prioritize, start with the first ones on the list for the week.
  - All readings available digitally on Canvas

# Course workflow

- **Problem sets** (25% of your grade)
  - Main day-to-day component of the class -- meant to get you working with your colleagues and thinking hard about the material.
  - Collaboration is **strongly encouraged** -- you should ask and answer questions on our class Ed Discussion board
  - Grading is on a **plus/check/minus** scale.
  - Conversion to grade point is somewhat hollistic.
  - Majority plusses is an A, Majority checks is an A-, Majority minus is a B
  - Solutions will be posted after the due-date.

# Course workflow

- **Midterm Exam** (30% of your grade)
  - The midterm will be structured like the problem sets with two main differences:
    1. You have about 1 week to complete them instead of 2
    2. You may **not** collaborate with one another on the exams.
- **Final Exam** (35% of your grade)
  - The final exam will be in-person and written.
  - Combination of theoretical questions and code/results analysis.
- **Participation** (10% of your grade)
  - It is important that you actively engage with lecture and section -- ask and answer questions.
  - Do the reading!
  - Participating on Ed counts towards this as well.

# Class Requirements

- **Overall:** An interest in learning and willingness to ask questions.
- Generally assume some background in probability theory and statistics (e.g. an intro course taught in most departments)
  - Regression a plus but not required -- if you know what  $(X'X)^{-1}X'Y$  is, then that is great, but not strictly necessary.
- Main concepts to be familiar with:
  - properties of random variables
  - estimands and estimators
  - bias, variance, consistency
  - central limit theorem
  - confidence intervals; hypothesis testing

# A brief overview

- **Week 1:** Introduction to potential outcomes
  - Defining causal estimands
- **Week 2-3:** Experiments
  - Why they work, design, inference
- **Week 4-5:** Selection-on-observables
  - Key assumptions, graphical models, how to think about confounding
  - Estimation via regression, weighting, matching + combinations
- **Week 6-8:** Designs for dealing with unobserved confounding
  - Instrumental variables
  - Differences-in-differences
  - Regression discontinuity
- **Week 9:** Causal mediation; Sensitivity analyses



# Random variables and estimators

# Random variables

- Understanding the behavior and properties of **random variables** is at the core of statistical theory.
- (Simply put) a random variable  $X$  is a mapping from a **sample space** to the real number line - Random variables have a *distribution* (which we may or may not assume we know) defined by the cumulative distribution function (CDF)

$$F(x) = \Pr(X \leq x)$$

# Random variables

- Discrete random variables take on a countable number of values (e.g. Bernoulli r.v. can take on 0 or 1) and have a probability mass function (PMF)

$$p(x) = \Pr(X = x)$$

- Continuous random variables take on an uncountable number of values (e.g. the Normal distribution on  $(-\infty, \infty)$ ). No PMF, but have a *density* function (PDF) that integrates to a probability

$$\Pr(X \in \mathcal{A}) = \int_{\mathcal{A}} f(x) dx$$

- **Remember:** PMFs (and PDFs) sum (integrate) to 1 over the support of the random variable.

# Expectations

- One important property of a random variable is its expectation  $E[X]$ . We'll often make assumptions about the expectation of an R.V. while remaining agnostic about its true distribution.
- The expectation is a *weighted average*. For a discrete r.v.  $X$ , we sum over the support of the random variable  $\mathcal{X}$ .

$$E[X] = \sum_{x \in \mathcal{X}} x \Pr(X = x)$$

- For continuous r.v. we have an integral

$$E[X] = \int_{x \in \mathcal{X}} x f(x) dx$$

- Fun fact: we can get the expectation of any function of  $g(X)$  just by plugging it into the integral

$$E[g(X)] = \int_{x \in \mathcal{X}} g(x) f(x) dx$$

# Expectations

- You'll probably not have to do any integration in this class. Rather, we'll derive expectations of functions of random variables via known properties
- Most important. **Linearity**. For any two random variables  $X$  and  $Y$  and constants  $a$  and  $b$

$$E[aX + bY] = aE[X] + bE[Y]$$

- Note that for any generic function  $g()$ ,  $E[g(X)] \neq g(E[X])$ . If  $g()$  is convex, by Jensen's inequality  $E[g(X)] \geq g(E[X])$
- For a binary r.v.  $X \in \{0, 1\}$ , it's helpful to remember the "fundamental bridge" between expectations and probability

$$E[X] = \Pr(X = 1)$$

# Variance

- We also care about the *spread* of a random variable -- how far is the average draw of  $X$  from its mean  $E[X]$ . One measure of this is the variance.

$$\text{Var}(X) = E[(X - E[X])^2]$$

- Also written as

$$\text{Var}(X) = E[X^2] - E[X]^2$$

- Note that the square is a convex function. Which means that by Jensen's inequality  $E[X^2] \geq E[X]^2$ . Variances *cannot* be negative!
- We also can define a *covariance* between two variables (does  $X$  take high values when  $Y$  takes high values?)

$$\text{Cov}(X, Y) = E[(X - E[X])(Y - E[Y])] = E[XY] - E[X]E[Y]$$

# Variance

- Again, you won't have to do the integral. Variances have some useful properties.
- For a constant  $a$

$$\text{Var}(aX) = a^2 \text{Var}(X)$$

- For any two random variables  $X$  and  $Y$

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y) + 2\text{Cov}(X, Y)$$

$$\text{Var}(X - Y) = \text{Var}(X) + \text{Var}(Y) - 2\text{Cov}(X, Y)$$

- For independent random variables  $X$  and  $Y$

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y)$$

$$\text{Var}(X - Y) = \text{Var}(X) + \text{Var}(Y)$$

# Conditional probabilities

- We will also spend *a lot of time* with conditional distributions and conditional expectations of random variables.
  - What's the probability that an individual enrolls in a job training program given their income?
  - We represent the conditioning set using a vertical bar with the right-hand side denoting what is being conditioned on.
    - For example:  $Pr(D_i = 1 | X_i = x)$



# Conditional probabilities

- **Key concept** - Dependence and independence. If two variables are *independent*, the distribution of one does not change conditional on the other. We'll write this using the  $\perp\!\!\!\perp$  notation.
- $Y_i \perp\!\!\!\perp D_i$  implies

$$Pr(Y_i = 1|D_i = 1) = Pr(Y_i = 1|D_i = 0) = Pr(Y_i = 1)$$

- Otherwise, the two variables are dependent:  $Y_i \not\perp\!\!\!\perp D_i$ .

$$Pr(Y_i = 1|D_i = 1) \neq Pr(Y_i = 1|D_i = 0)$$

- Two variables can be *conditionally independent* in that they are independent only when conditioning on a third variable. For example, we can have  $Y_i \perp\!\!\!\perp D_i$  but  $Y_i \not\perp\!\!\!\perp D_i|X_i$ . This implies

$$Pr(Y_i = 1|D_i = 1, X_i = x) = Pr(Y_i = 1|D_i = 0, X_i = x) = Pr(Y_i = 1|X_i = x)$$

- **Remember:** Conditional independence *does not* imply independence or vice-versa!

# Conditional expectations

- A central object of interest in statistics is the **conditional expectation** function (CEF)  $E[Y|X]$ .
  - Given a particular value of  $X$ , what is the expectation of  $Y$ ?
  - The CEF is a function of  $X$ .
- All the usual properties of expectations apply to conditional expectations.
- We also will often make use of the *law of total expectation*

$$E[Y] = E[E[Y|X]]$$

- Easiest to think about this in terms of discrete r.v.s

$$E[Y] = \sum_{x \in \mathcal{X}} E[Y|X = x] Pr(X = x)$$

# Estimation

- One critical use of statistical theory is understanding how to learn about things we **don't observe** using things that we **do observe**. We call this estimation.
  - e.g. What is the share of voters in Wisconsin who will turn out in the 2022 election?
  - What is the share of voters who turn out among those assigned to receive a GOTV phone call?
- **Estimand**: The **unobserved** quantity that we want to learn about. Often denoted via a greek letter (e.g.  $\mu$ ,  $\pi$ )
  - Often a "population" characteristic that we want to learn about via a sample.
    - But you'll learn another reason why we sometimes can't observe a quantity of interest even in a sample!
  - Important to define your estimand well. (Lundberg, Johnson and Stewart, 2022)

# Estimation

- **Estimator**: The **function** of random variables that we will use to try to estimate the quantity of interest. Often denoted with a hat on the parameter of interest (e.g.  $\hat{\mu}$ ,  $\hat{\pi}$ )
  - Why are the variables random?
    - Classic inference: We have a random sample from the population -- if we took another sample, we would obtain a different realization of our estimator.
    - Randomization inference: We have a randomly assigned treatment -- if we were to re-run the experiment, we would observe a different treatment/control allocation.
- **Estimate**: A single realization of our estimator (e.g. 0.3, 9.535)
  - We often report both point estimates ("best guess") and interval estimates (e.g. confidence intervals).
  - Careful not to confuse properties of estimators with properties of the estimates themselves.

# Estimation



estimand

Ingredients	Method
150g unsalted butter, plus extra for greasing	1. Heat the oven to 160C/140C fan/gas 3. Grease and base line a 1 litre heatproof glass pudding basin and a 450g loaf tin with baking parchment.
150g plain chocolate, broken into pieces	
150g plain flour	
½ tsp baking powder	2. Put the butter and chocolate into a saucepan and melt over a low heat, stirring. When the chocolate has all melted remove from the heat.
½ tsp bicarbonate of soda	
200g light muscovado sugar	
2 large eggs	

estimator



estimate

# Estimation

- The classic estimation problem in statistics is to estimate some unknown population mean  $\mu$  from an i.i.d. sample of  $n$  observations  $Y_1, Y_2, \dots, Y_n$ .
  - We assume that each  $Y_i$  is a draw from the target population with mean  $\mu$ . (identically distributed) -- therefore  $E[Y_i] = \mu$
  - We'll also assume that conditioning on  $Y_i$  tells us nothing about any other  $Y_j$   $Y_i \perp\!\!\!\perp Y_j$  (independently distributed) -- this implies  $Cov(Y_i, Y_j) = 0$
- Our estimand:  $\mu$
- Our estimator: The sample mean  $\hat{\mu} = \bar{Y} = \frac{1}{n} \sum_{i=1}^n Y_i$
- Our estimate: A particular realization of that estimator based on our observed sample (e.g. 0.4)
- Note that our estimator is a *random variable* -- it's a function of  $Y_i$ s which are random variables.
  - Therefore it has an expectation  $E[\hat{\mu}]$  (assuming  $Y_i$  has an expectation)
  - It has a variance  $Var(\hat{\mu})$  (again, under regularity conditions)
  - It has a distribution (which we may or may not know).

# Estimation

- How do we know if we've picked a good estimator? Will it be close to the truth? Will it be systematically higher or lower than the target?
- We want to derive some of its properties
  - Bias:  $E[\hat{\mu}] - \mu$
  - Variance:  $Var(\hat{\mu})$
  - Consistency: Does  $\hat{\mu}$  converge in probability to  $\mu$  as  $n$  goes to infinity?
  - Asymptotic distribution: Is the sampling distribution of  $\hat{\mu}$  well approximated by a known distribution?

# Unbiasedness

- Is the expectation of  $\hat{\mu}$  equal to  $\mu$ ?
- First we pull out the constant.

$$E[\hat{\mu}] = E\left[\frac{1}{n} \sum_{i=1}^n Y_i\right] = \frac{1}{n} E\left[\sum_{i=1}^n Y_i\right]$$

- Next we use linearity of expectations

$$\frac{1}{n} E\left[\sum_{i=1}^n Y_i\right] = \frac{1}{n} \sum_{i=1}^n E[Y_i]$$

- Finally, under our i.i.d. assumption

$$\frac{1}{n} \sum_{i=1}^n E[Y_i] = \frac{1}{n} \sum_{i=1}^n \mu = \frac{n\mu}{n} = \mu$$

- Therefore, the bias,  $\text{Bias}(\hat{\mu}) = E[\hat{\mu}] - \mu = 0$



# Variance

- What is the variance of  $\hat{\mu}$ ? Again, start by pulling out the constant.

$$\text{Var}(\hat{\mu}) = \text{Var} \left[ \frac{1}{n} \sum_{i=1}^n Y_i \right] = \frac{1}{n^2} \text{Var} \left[ \sum_{i=1}^n Y_i \right]$$

- We can further simplify by using our i.i.d. assumption. The variance of a sum of i.i.d. random variables is the sum of the variances.

$$\frac{1}{n^2} \text{Var} \left[ \sum_{i=1}^n Y_i \right] = \frac{1}{n^2} \sum_{i=1}^n \text{Var} [Y_i]$$

- "identically distributed"

$$\frac{1}{n^2} \sum_{i=1}^n \text{Var} [Y_i] = \frac{1}{n^2} \sum_{i=1}^n \sigma^2 = \frac{n\sigma^2}{n^2} = \frac{\sigma^2}{n}$$

- Therefore, the variance is  $\frac{\sigma^2}{n}$

# Asymptotic behavior

- As  $n$  gets large, what can we say about the estimator  $\hat{\mu}$ .
- First, we can show that it is **consistent** -- it converges in probability to the true parameter  $\mu$ 
  - Unbiasedness + Variance that goes to 0 as  $n$  gets large.
  - Some estimators may be biased but have bias terms that go to 0 -- if variance also goes to 0 these are still consistent.
- Second, we can say something about the distribution of  $\hat{\mu}$ .
  - Remember, we've only made assumptions about  $E[Y_i]$  and  $Var(Y_i)$  (that they exist). We have made no assumptions on the distribution of  $Y_i$ .  $Y_i$  can be normal, poisson, bernoulli, or whatever!
  - However, we know something about sums and means of random variables -- they are well-approximated by a normal distribution. **The Central Limit Theorem!**
  - So in large samples, the **sampling distribution** of  $\hat{\mu}$  is close to normal. This lets us construct confidence intervals and do inference with this approximation and be confident that we won't be far off!

The potential outcomes model

# Thinking about causal effects

- Two types of causal questions (Gelman and Rubin, 2013)
- **Causes of effects**
  - What are the factors that cause some outcome  $Y$ ?
  - "Why?" questions: Why do states go to war? Why do politicians get re-elected?
- **Effects of causes**
  - If  $X$  were to change, what might happen to  $Y$ ?
  - "What if?" questions: What if a pair of states were democratic, would that change their chances of going to war? What if a politician were an incumbent, would that affect their re-election probability?
- We'll spend more time on the **effects of causes**
  - Why? Because we can connect them to well-defined statistical quantities of interest (e.g. an "average treatment effect")
  - "Causes of effects" are still important questions, and theoretical inquiry can lead you towards possible causes to evaluate.

# Defining a causal effect

- Historically, causality was seen as a *deterministic* process.
  - Hume (1740): Causes are regularities in events of "constant conjunctions"
  - Mill (1843): Method of difference
- This became problematic -- empirical observation alone does not demonstrate causality.
  - Russell (1913): Scientists aren't interested in causality!
- How do we talk about causation that both incorporates *uncertainty* in measurement and clearly defines what we mean by a "causal effect"?
- **Rubin** (1974) - formalizes a framework for understanding causation from a statistical perspective. Inspired by earlier Neyman (1923) and Fisher (1935) on randomized experiments.
- We'll spend most of our time with this approach, often called the **Rubin Causal Model** or **potential outcomes** framework.

# Causality and interventions

- Causal effects are effects of **interventions**
  - What happens to an outcome when a treatment is changed.
- It's very difficult to learn about vague causal statements:
  - What is the effect of growth on democracy or conflict?
  - What is the effect of exercise on health?
- The **potential outcomes** framework clarifies:
  1. **What** action is doing the causing?
  2. Compared to **what** alternative action?
  3. On **what** outcome metric?
  4. How would we learn about the effect from data?

# Statistical setup.

- Population of units
  - Finite population or infinite super-population
- Sample of  $N$  units from the population indexed by  $i$  - Observed outcome  $Y_i$
- Binary treatment indicator  $D_i$ .
  - Units receiving "treatment":  $D_i = 1$
  - Units receiving "control":  $D_i = 0$
- Covariates (observed prior to treatment)  $X_i$

# Potential outcomes

- Let  $D_i$  be the value of a **treatment** assigned to each individual.
- $Y_i(d)$  is the value that the outcome would take if  $D_i$  were set to  $d$ .
  - For binary  $D_i$ :  $Y_i(1)$  is the value we would observe if unit  $i$  were *treated*.
  - $Y_i(0)$  is the value we would observe if unit  $i$  were under *control*
- We model potential outcomes as **fixed** attributes of the units.
- **Notation alert!** -- Sometimes you'll see potential outcomes written as:
  - $Y_i^1, Y_i^0$  or  $Y_i^{d=1}, Y_i^{d=0}$
  - $Y_{i0}, Y_{i1}$
  - $Y_1(i), Y_0(i)$
- Causal effects are contrasts in potential outcomes.
  - Individual treatment effect:  $\tau_i = Y_i(1) - Y_i(0)$
  - Can consider ratios or other transformations (e.g.  $\frac{Y_i(1)}{Y_i(0)}$ )



# Consistency/SUTVA

- How do we link the potential outcomes to observed ones?
- Consistency/Stable Unit Treatment Value (SUTVA) assumption

$$Y_i(d) = Y_i \text{ if } D_i = d$$

- Sometimes you'll see this w/ binary  $D_i$  (often in econometrics)

$$Y_i = Y_i(1)D_i + Y_i(0)(1 - D_i)$$

- Implications
  1. No interference -- other units' treatments don't affect  $i$ 's potential outcomes.
  2. Single version of treatment
  3.  $D$  is in principle manipulable -- a "well-defined intervention"
  4. The means by which treatment is assigned is irrelevant (a version of 2)

# Positivity/Overlap

- We also need some assumptions on the treatment assignment mechanism  $D_i$ .
- In order to be able to observe *some units'* values of  $Y_i(1)$  or  $Y_i(0)$  treatment can't be deterministic. For all  $i$ :

$$0 < \Pr(D_i = 1) < 1$$

- If no units could ever receive treatment or control it would be impossible to learn about  $E[Y_i|D_i = 1]$  or  $E[Y_i|D_i = 0]$
- This is sometimes called a **positivity** or overlap assumption.
  - Pretty trivial in a randomized experiment, but can be tricky in observational studies when  $D_i$  is perfectly determined by some covariates  $X_i$

# A missing data problem

- It's useful to think of the causal inference problem in terms of *missingness* in the complete table of potential outcomes.

Unit $i$	Treatment $D_i$	$Y_i(1)$	$Y_i(0)$	Observed $Y_i$
1	1	5	?	5
2	0	?	-3	-3
3	1	9	?	9
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$
$N$	0	?	8	8

- If we could observe both  $Y_i(1)$  and  $Y_i(0)$  for each unit, then this would be easy!
- But we can't - we only observe what we're given by  $D_i$
- [Holland \(1986\)](#) calls this "The Fundamental Problem of Causal Inference"

# Causal Estimands

- All causal inference starts with a definition of the estimand.
- The individual causal effect:  $\tau_i$

$$\tau_i = Y_i(1) - Y_i(0)$$

- **Problem:** Can't identify this without extremely strong assumptions!
  - "The Fundamental Problem of Causal Inference"
- The **average treatment effect** (ATE):  $\tau$

$$\tau = E[Y_i(1) - Y_i(0)]$$

- What is the *expected* effect of the treatment?

# Causal Estimands

- The **conditional average treatment effect** (CATE):  $\tau(x)$

$$\tau(x) = E[Y_i(1) - Y_i(0) | X_i = x]$$

- The average treatment effect for the sub-population with a particular set of covariate values  $X_i = x$

- The **average treatment effect on the treated** (ATT):  $\tau_{\text{ATT}}$

$$\tau_{\text{ATT}} = E[Y_i(1) - Y_i(0) | D_i = 1]$$

- The average treatment effect among those units that were *assigned* treatment.
- Why do we care about this?
  - Sometimes more policy-relevant.
  - Often easier to identify in selection-on-observables designs.

# Causal Estimands

- The **risk ratio** (RR) for binary outcomes:

$$RR = \frac{Pr(Y_i(1) = 1)}{Pr(Y_i(0) = 1)}$$

- Political science and social sciences rarely use these.
- Medicine *loves* risk ratios (e.g. Vaccine effectiveness = 1 - RR)

# Finite Population Estimands

- For these estimands we have two sources of uncertainty
  - Sampling from a population.
  - Random assignment of treatment (unobserved P.O.s)
- What if we're just interested in doing inference on the treatment effect in the sample, or there is no population from which we sample from?
  - We can define analogues of the ATE for the sample -- the *sample average treatment effect*

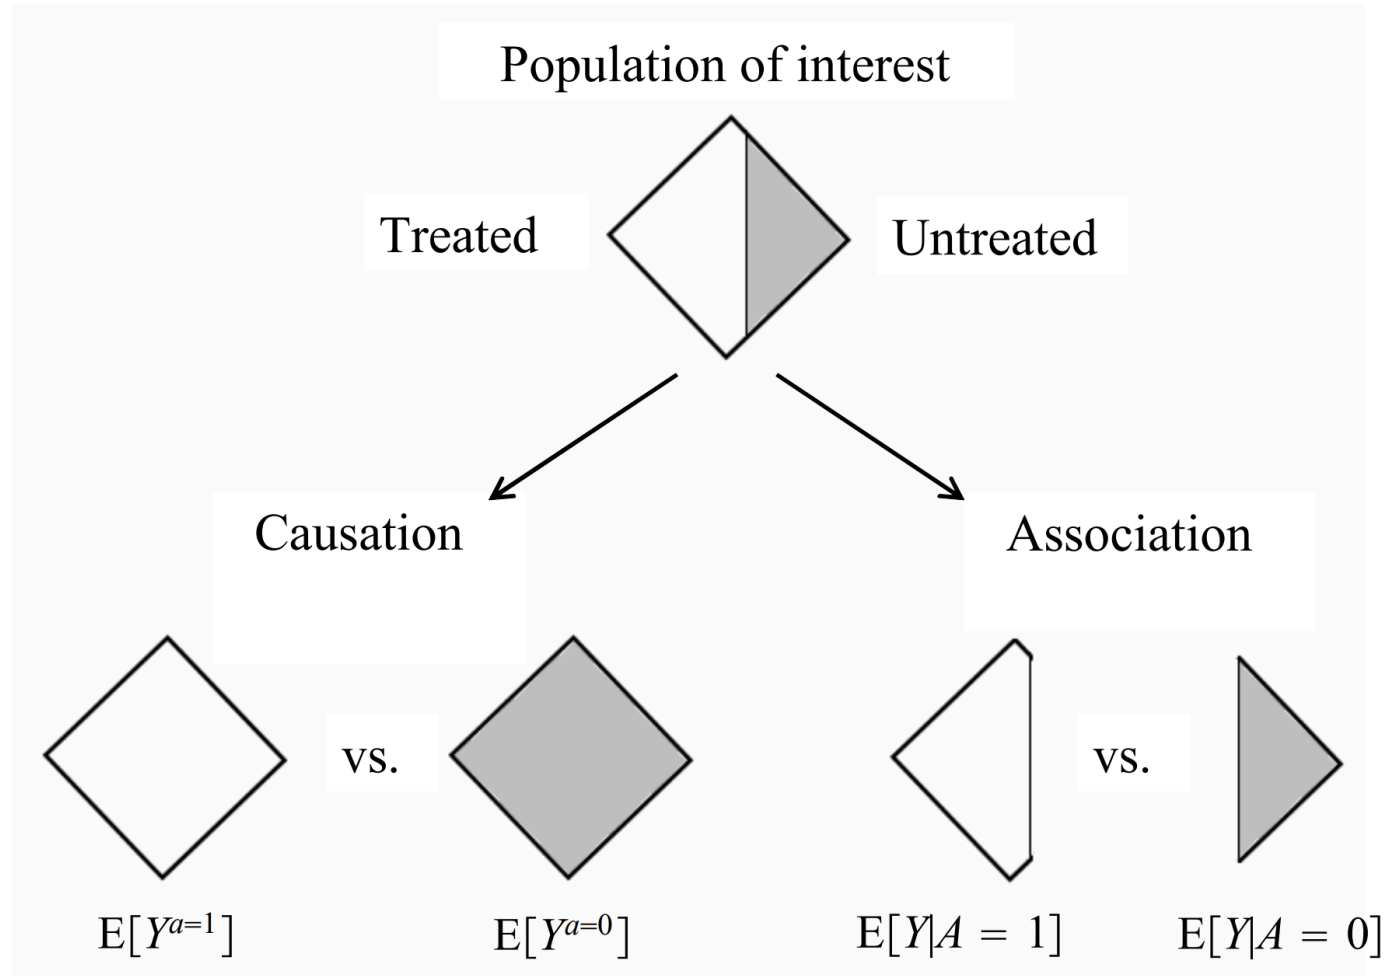
$$\tau_{\text{SATE}} = \frac{1}{N} \sum_{i=1}^N Y_i(1) - Y_i(0)$$

# Finite Population Estimands

- There's *still* uncertainty -- we don't observe all of the potential outcomes. But now we don't have to justify generalizing to the population.
  - Under random sampling, the SATE is unbiased for the (population) ATE or PATE.
  - Large literature on when and how we can generalize absent random sampling ("transportability")
- Inference justified on the basis of *randomization* of the treatment.
  - A lot of the same tools for inference in classical super-population sampling transfer over nicely.



# Causal vs. Associational Estimands



# Identification

- In statistics, we often talk about whether a parameter is *identifiable*.
  - A parameter is "point identified" if, having access to infinite data, the parameter could only take on a single value.
  - In other words, no other value of the parameter could generate the same observable data.
- In statistical models, non-identifiability arises when different values of the parameter could give rise to the same observable data
  - In classical regression: More parameters than observations.
  - Or perfectly collinear regressors

# Causal Identification

- **Causal identification:** Can we learn about the value of a causal effect from the observed data.
  - Can we express the causal estimand (e.g.  $\tau = E[Y_i(1) - Y_i(0)]$ ) entirely in terms of *observable* quantities
- Causal identification comes prior to questions of estimation
  - It doesn't matter whether you're using regression, weighting, matching, doubly-robust estimation, double-LASSO, etc...
  - If you can't answer the question "What's your identification strategy?" then no amount of fancy stats will solve your problems.
- Identification requires *assumptions* about the connection between the observed data  $Y_i$ ,  $D_i$  and the unobserved counterfactuals  $Y_i(d)$ 
  - Under what assumptions will the observed difference-in-means *identify* the average treatment effect?

# Conclusion

- Causal effects are contrasts in **counterfactuals**
- The potential outcomes framework gives us a tool for defining statistical *estimands* in terms of functions of individual effects of treatment.
  - The potential outcome  $Y_i(d)$  denotes the outcome  $Y_i$  we would observe for unit  $i$  if they were assigned treatment  $d$
  - Implicitly making a consistency/SUTVA assumption to link *observed* to *counterfactual*
- **Fundamental problem of causal inference**: We're going to need more assumptions -- don't get to directly observe individual treatment effects.
- Consistency and positivity only get us part of the way there.
- Next week: What additional assumptions do we need to **identify** causal effects from the observed data. Why randomized experiments satisfy those assumptions.