

Week 9: Sensitivity Analyses

PLSC 30600 - Causal Inference

Last three weeks

- Identification under *unobserved* confounding.
- **Solutions:** Make additional assumptions on the form of the confounding
- **Instrumental Variables**
 - Assume unobserved confounding doesn't affect the instrument.
- **Difference-in-differences**
 - Assume unobserved confounding affects "placebo" outcome the same way as the outcome of interest.
- **Regression discontinuity design**
 - No unobserved confounding in vicinity of a treatment assignment discontinuity

Today

- Back to **selection on observables**!
- How can we diagnose our no unobserved confounders assumption?
 - What sorts of confounders would threaten our results (cause them to go to 0).
- Omitted variable bias formula
 - Allows us to "sign the bias" of a proposed confounder based on the outcome-confounder and outcome-treatment relationships
- Sensitivity analysis
 - How bad of a hypothetical confounder would we need to break the result?
 - Provides a benchmark for any critiques of a selection-on-observables assumption.

Omitted variable bias

Omitted Variable Bias

- Let's return to our previous setting with treatment D_i and outcome Y_i .
- Suppose there exists an omitted confounder U_i and ignorability holds conditional on that omitted confounder.
 - Suppose we ignore it and just use a simple difference-in-means estimator.
- What's the bias for the ATT? Recall our selection-into-treatment bias formula!

$$\underbrace{E[Y_i|D_i = 1] - E[Y_i|D_i = 0]}_{\text{Difference-in-means}} = \underbrace{E[Y_i(1) - Y_i(0)|D_i = 1]}_{\text{ATT}} + \left(\underbrace{E[Y_i(0)|D_i = 1] - E[Y_i(0)|D_i = 0]}_{\text{Selection-into-treatment bias}} \right)$$

Omitted Variable Bias

- Let's write the selection bias conditioning on U_i

$$\text{Selection Bias} = \sum_{u \in \mathcal{U}} E[Y_i(0) | D_i = 1, U_i = u] Pr(U_i = u | D_i = 1) - \sum_{u \in \mathcal{U}} E[Y_i(0) | D_i = 0, U_i = u] Pr(U_i = u | D_i = 0)$$

- Ignorability conditional on U_i

$$\text{Selection Bias} = \sum_{u \in \mathcal{U}} E[Y_i(0) | U_i = u] Pr(U_i = u | D_i = 1) - \sum_{u \in \mathcal{U}} E[Y_i(0) | U_i = u] Pr(U_i = u | D_i = 0)$$

- Combining terms

$$\text{Selection Bias} = \sum_{u \in \mathcal{U}} E[Y_i(0) | U_i = u] \times \left(Pr(U_i = u | D_i = 1) - Pr(U_i = u | D_i = 0) \right)$$

Omitted Variable Bias

- Two elements to selection bias. First, if treatment assignment is independent of the confounder, then the bias is 0

$$\text{Selection Bias} = \sum_{u \in \mathcal{U}} E[Y_i(0) | U_i = u] \times \left(\Pr(U_i = u | D_i = 1) - \Pr(U_i = u | D_i = 0) \right)$$

- Second, if $Y_i(0)$ is independent of U_i , we have:

$$\text{Selection Bias} = \sum_{u \in \mathcal{U}} E[Y_i(0)] \times \left(\Pr(U_i = u | D_i = 1) - \Pr(U_i = u | D_i = 0) \right)$$

$$\text{Selection Bias} = E[Y_i(0)] \times \left(\sum_{u \in \mathcal{U}} \Pr(U_i = u | D_i = 1) - \sum_{u \in \mathcal{U}} \Pr(U_i = u | D_i = 0) \right)$$

$$\text{Selection Bias} = E[Y_i(0)] \times (1 - 1) = 0$$

- We get OVB/confounding when:
 - U_i is not independent of treatment
 - U_i is not independent of the potential outcomes

Signing the bias

- Additionally, the bias is multiplicative.
- Under some constant effects assumptions, we can get the direction of the bias of the difference-in-means relative to the ATT
 1. **Positive** association between the confounder on outcome. **Positive** association between confounder and treatment. **Positive** bias.
 2. **Positive** association between the confounder on outcome. **Negative** association between confounder and treatment. **Negative** bias.
 3. **Negative** association between the confounder on outcome. **Positive** association between confounder and treatment. **Negative** bias.
 4. **Negative** association between the confounder on outcome. **Negative** association between confounder and treatment. **Positive** bias.

Example: Smoking and Cancer

- Back when the link between smoking and cancer was being debated, some researchers suggested that cigarettes might be a "healthy" alternative to pipe smoking
- Cochran (1968) uses this to illustrate adjustment by stratification

TABLE 1
DEATH RATES PER 1,000 PERSON-YEARS

Smoking group	Study		
	Canadian	British	U. S.
Non-smokers	20.2	11.3	13.5
Cigarettes only	20.5	14.1	13.5
Cigars, pipes	35.5	20.7	17.4

1. What's the omitted confounder?
2. What's the direction of the bias due to the omitted confounder?

Example: Smoking and Cancer

TABLE 2
MEAN AGES, YEARS

Smoking group	Study		
	Canadian	British	U. S.
Non-smokers	54.9	49.1	57.0
Cigarettes only	50.5	49.8	53.2
Cigars and/or pipe	65.9	55.7	59.7

Example: Smoking and Cancer

TABLE 3

ADJUSTED DEATH RATES USING 2, 3, AND 9-11 SUBCLASSES

Number of subclasses	Canadian			British			U. S.		
	N. S.*	C.+	CP'	N. S.	C	CP	N. S.	C	CP
1	20.2	20.5	35.5	11.3	14.1	20.7	13.5	13.5	17.4
2	20.2	26.4	24.0	11.3	12.7	13.6	13.5	16.4	14.9
3	20.2	28.3	21.2	11.3	12.8	12.0	13.5	17.7	14.2
9-11	20.2	29.5	19.8	11.3	14.8	11.0	13.5	21.2	13.7

*Non-smokers, +Cigarettes only, 'Cigars, Pipes

OVB in linear models

- Suppose we want to identify the effect of D on Y conditional on pre-treatment covariates \mathbf{X} . Assume we're willing to assume a linear model for the outcome and that there exists one omitted covariate Z

$$Y = \hat{\tau}D + \mathbf{X}\hat{\beta} + \hat{\gamma}Z + \hat{\epsilon}$$

- What happens if we instead estimate the "restricted" model with Z omitted?

$$Y = \hat{\tau}_{\text{res}}D + \mathbf{X}\hat{\beta}_{\text{res}} + \hat{\epsilon}_{\text{res}}$$

- What's the relationship between $\hat{\tau}_{\text{res}}$?

OVB in linear models

- Let's define $D^{\perp \mathbf{X}}$ as the "partialled-out" value of D (the residuals from a regression of D on X).
 - Similarly define $Y^{\perp \mathbf{X}}$ as the "partialled-out" value of Y given X
- By the Frisch-Waugh-Lovell theorem, we can write any regression coefficient in terms of the "partialled" bivariate regression

$$\hat{\tau}_{\text{res}} = \frac{\text{cov}(D^{\perp \mathbf{X}}, Y^{\perp \mathbf{X}})}{\text{var}(D^{\perp \mathbf{X}})}$$

- Using our definition of Y (by the linear model)

$$\hat{\tau}_{\text{res}} = \frac{\text{cov}(D^{\perp \mathbf{X}}, \hat{\tau} D^{\perp \mathbf{X}} + \hat{\gamma} Z^{\perp \mathbf{X}})}{\text{var}(D^{\perp \mathbf{X}})}$$

- Properties of covariance

$$\hat{\tau}_{\text{res}} = \hat{\tau} \frac{\text{cov}(D^{\perp \mathbf{X}}, D^{\perp \mathbf{X}})}{\text{var}(D^{\perp \mathbf{X}})} + \hat{\gamma} \frac{\text{cov}(D^{\perp \mathbf{X}}, Z^{\perp \mathbf{X}})}{\text{var}(D^{\perp \mathbf{X}})}$$

OV B in linear models

- Simplifying

$$\hat{\tau}_{\text{res}} = \hat{\tau} + \hat{\gamma} \frac{\text{cov}(D^{\perp \mathbf{X}}, Z^{\perp \mathbf{X}})}{\text{var}(D^{\perp \mathbf{X}})}$$

- We can recognize that the last term is the coefficient on D from a regression of Z on D and X (again using FWL) - let's call this $\hat{\delta}$

$$\hat{\tau}_{\text{res}} = \hat{\tau} + \hat{\gamma} \hat{\delta}$$

- So the discrepancy between the "restricted" and "unrestricted" models can be written as the product of two coefficients - the relationship between Z and Y (given X) and the relationship between Z and D (given D)

$$\widehat{\text{bias}} = \hat{\gamma} \hat{\delta}$$

Signing the bias

- Gentzkow, Shapiro, Sinkinson (2011: AER) examine the effect of newspaper entry on political competitiveness in the US counties from 1869 to 1928.
 - Outcome: Presidential/congressional turnout
 - Treatment: Number of new newspapers
 - Finding: More newspapers increase turnout.
- Consider some hypothetical confounders, in what direction would they bias the estimate?
 - **Population growth**: How would population growth affect newspaper entry? How would it affect turnout?
 - **Income growth** How would income growth affect newspaper entry and turnout?
- How would we expect either of these confounders to alter our estimate?

Sensitivity analyses

Sensitivity analysis

- Sensitivity analyses ask the question "how bad of a violation of our identification assumptions would break our result?"
 - We use a parameter (or parameters) to represent the violation and re-do the analysis.
 - Vary the parameter over a (sensible) range - how often do our results appreciably change (e.g. become zero or flip sign)
- **Challenge:** How do we define a suitable sensitivity parameter that has actual interpretability?

Confounding function

- A general approach to thinking about sensitivity parameters in a binary treatment setting comes from **Blackwell (2014)**
- Define the "confounding function"

$$q(d, x) = E[Y_i(d)|D_i = d, X_i = x] - E[Y_i(d)|D_i = 1 - d, X_i = x]$$

- The confounding function captures the extent to which the potential outcomes **differ** between a treated unit and a control unit with $X_i = x$.
 - Under ignorability, $q(d, x) = 0$
- We could set the confounding function to have a particular form:
 - For example, $q(d, x) = \alpha$ implies the selection bias is **constant** at all levels of X_i

Sensitivity analyses

- Given a value of $q(d, x)$, we can straightforwardly implement a sensitivity analysis by **de-biasing** the outcome

$$Y_i^q = Y_i - q(D_i, X_i) \times Pr(1 - D_i | X_i)$$

- Then, run the analysis on Y_i^q
 - Vary the sensitivity parameters for α and see what magnitude of confounding is enough to "break" the results.

Sensitivity analyses

- The intuition for the debiasing comes from our selection-into-treatment bias decomposition.
Without covariates:

$$\begin{aligned} E[Y_i(0)] &= E[Y_i(0)|D_i = 0]Pr(D_i = 0) + E[Y_i(0)|D_i = 1]Pr(D_i = 1) \\ &= E[Y_i(0)|D_i = 0] - \left(E[Y_i(0)|D_i = 0] - E[Y_i(0)|D_i = 1] \right) \times Pr(D_i = 1) \\ &= E[Y_i|D_i = 0] - q(0) \times Pr(D_i = 1) \\ &= E[Y_i^q|D_i = 0] \end{aligned}$$

Sensitivity analysis in linear models

- An alternative approach is to think about confounding in terms of **two** quantities
 - The relationship between treatment and confounder
 - The relationship between outcome and confounder
- In a linear model setting, we could construct a sensitivity analysis in terms of two parameters:
 - The partial regression coefficient between Z and Y , $\hat{\gamma}$
 - The partial regression coefficient between Z and D , $\hat{\delta}$
- Slightly annoying since each of these depends on the scale of Z and Y - can we re-write in terms of parameters with the same range irrespective of the outcome?
 - Cinelli and Hazlett (2020) provide a reparameterization in terms of the partial R^2 of two regressions involving Z (which are always between 0 and 1)

Rewriting the bias

- Start by defining the $R^2_{Z \sim D}$ as the R^2 from a regression of Z on D .
 - For OLS: $R^2_{Z \sim D} = 1 - \frac{\text{Var}(Z^{\perp D})}{\text{Var}(Z)} = \text{cor}(Z, D)^2 = \left(\frac{\text{cov}(Z, D)}{\text{sd}(Z)\text{sd}(D)} \right)^2$
 - Same thing for the partial R^2 : $R^2_{Z \sim D | \mathbf{X}} = \text{cor}(Z^{\perp \mathbf{X}}, D^{\perp \mathbf{X}})^2$
- Now write our bias term

$$\widehat{\text{bias}} = \left(\frac{\text{cov}(D^{\perp \mathbf{X}}, Z^{\perp \mathbf{X}})}{\text{var}(D^{\perp \mathbf{X}})} \right) \left(\frac{\text{cov}(Y^{\perp \mathbf{X}, D}, Z^{\perp \mathbf{X}, D})}{\text{var}(Z^{\perp \mathbf{X}, D})} \right)$$

- Convert covariance to correlation

$$\widehat{\text{bias}} = \left(\frac{\text{cor}(D^{\perp \mathbf{X}}, Z^{\perp \mathbf{X}})\text{sd}(Z^{\perp \mathbf{X}})}{\text{sd}(D^{\perp \mathbf{X}})} \right) \left(\frac{\text{cor}(Y^{\perp \mathbf{X}, D}, Z^{\perp \mathbf{X}, D})\text{sd}(Y^{\perp \mathbf{X}, D})}{\text{sd}(Z^{\perp \mathbf{X}, D})} \right)$$

Rewriting the bias

- Rearrange terms

$$\widehat{\text{bias}} = \left(\frac{\text{cor}(D^{\perp \mathbf{X}}, Z^{\perp \mathbf{X}}) \text{cor}(Y^{\perp \mathbf{X}, D}, Z^{\perp \mathbf{X}, D})}{\frac{\text{sd}(Z^{\perp \mathbf{X}, D})}{\text{sd}(Z^{\perp \mathbf{X}})}} \right) \left(\frac{\text{sd}(Y^{\perp \mathbf{X}, D})}{\text{sd}(D^{\perp \mathbf{X}})} \right)$$

- Square everything

$$\widehat{\text{bias}}^2 = \left(\frac{\text{cor}(D^{\perp \mathbf{X}}, Z^{\perp \mathbf{X}})^2 \text{cor}(Y^{\perp \mathbf{X}, D}, Z^{\perp \mathbf{X}, D})^2}{\frac{\text{var}(Z^{\perp \mathbf{X}, D})}{\text{var}(Z^{\perp \mathbf{X}})}} \right) \left(\frac{\text{var}(Y^{\perp \mathbf{X}, D})}{\text{var}(D^{\perp \mathbf{X}})} \right)$$

- Substitute in the (partial) R^2 parameters

$$\widehat{\text{bias}}^2 = \left(\frac{R_{D \sim Z | \mathbf{X}}^2 R_{Y \sim Z | \mathbf{X}, D}^2}{1 - R_{D \sim Z | \mathbf{X}}^2} \right) \left(\frac{\text{var}(Y^{\perp \mathbf{X}, D})}{\text{var}(D^{\perp \mathbf{X}})} \right)$$

Rewriting the bias

- Take the square root to get the absolute bias

$$|\widehat{\text{bias}}| = \sqrt{\frac{R_{D \sim Z|X}^2 R_{Y \sim Z|X,D}^2}{1 - R_{D \sim Z|X}^2}} \left(\frac{\text{sd}(Y^{\perp X,D})}{\text{sd}(D^{\perp X})} \right)$$

- Some important intuitions
 - The bias is a *product* of the magnitude of the two R^2 . An unobserved confounder that explains very little of the treatment needs to explain *a lot* of the outcome to induce a sizeable bias (and vice-versa)
 - The bias is smaller when the amount of variation in the outcome given X and D is low (not much Y left to explain)
 - The bias is *amplified* when the variability in D given X is low.

Illustration: Hazlett (2019)

- Hazlett (2019; JCR) examines the impact of exposure to violence on attitudes towards peace in the context of the war in Darfur.
 - Key finding: Refugees with greater exposure to violence are more likely to express support for peace - support for a "war-weariness" theory of attitudes during conflict.
 - Identification strategy: Selection-on-observables conditional on village and gender (plus other covariates).
 - Argues that exposure to violence by pro-government militias across villages was non-random but within-village often indiscriminate.
- How bad does the residual confounding need to be to break the result?

Illustration: Hazlett (2019)

```
library(sensemakr)
data('darfur')

darfur.reg <- lm(peacefactor ~ directlyharmed + village + female + age + farmer_dar + herder_da
tidy(darfur.reg) %>% filter(term == "directlyharmed")
```

```
## # A tibble: 1 × 5
##   term          estimate std.error statistic   p.value
##   <chr>         <dbl>    <dbl>    <dbl>   <dbl>
## 1 directlyharmed  0.0973    0.0233     4.18 0.0000318
```

```
sd(darfur$peacefactor)
```

```
## [1] 0.348
```

Illustration: Hazlett (2019)

- We would like to generate a plot of how the results would change as we vary the two R^2 parameters by calculating the bias across each of the possible parameter values.
 - Need 3 dimensions (each R^2 plus the estimate)
 - Can do this manually...but extremely tedious - luckily Cinelli and Hazlett make a great R package **sensemakr**

```
sensitivity <- sensemakr(model = darfur.reg,  
                        treatment = "directlyharmed",  
                        benchmark_covariates = "female",  
                        kd = 1:3,  
                        ky = 1:3,  
                        q = 1,  
                        alpha = .05,  
                        reduce = T)
```

Illustration: Hazlett (2019)




Illustration: Hazlett (2019)

- "Robustness Value" - consider a confounder that has an equal partial R^2 with the outcome and the treatment -- what's the smallest such R^2 necessary to drive the result to 0 (or insignificance).

```
sensitivity$sensitivity_stats
```

```
##          treatment estimate      se t_statistic r2yd.x  rv_q  rv_qa f2yd.x dof
## 1 directlyharmed    0.0973 0.0233      4.18 0.0219 0.139 0.0763 0.0224 783
```

- While the R^2 stats do have an intuitive interpretation, there is no *absolute* scale for what constitutes a "robust" vs. "non-robust" result.
 - It depends *also* on how much unexplained variability there is in the outcome.
 - Useful to *benchmark* the unobserved confounding against other known confounders

Illustration: Hazlett (2019)

- Consider the previous contour plot - the red points indicate bias under hypothetical confounding that is $1x$, $2x$, and $3x$ as strong as gender
- Why pick gender (versus any other observed confounder)?
 - We have prior theoretical reasons to believe it's strongly associated with both outcome and treatment.
- Caution with *informal* benchmarking - Cinelli and Hazlett (2020) show that just calculating the observed partial R^2 s for the benchmarks can be inaccurate.
 - Estimates of how X relates to Y may be biased due to omission of Z .
 - Also D is a collider.
- C+H (2020) derive formal bounds for the "benchmark" exercise with a set of observed covariates.

Summary

- Sensitivity analyses are a tool for *arguing*
 - You can always find values of the sensitivity parameters for which the results fail to hold
 - You can always find values of the sensitivity parameters for which the results *do* hold.
- What sensitivity analyses do is describe how *severe* a violation of selection-on-observables needs to be in order to threaten the main conclusions.
 - (e.g.) Smoking and cancer -- even if there were some unobserved confounder it would need to explain an *enormous* amount of variance for us to conclude no effect.
- Norms are developing about sensitivity analyses and reporting in observational designs
 - "Robustness Values" and contour plots
 - How best to "benchmark" -- what to benchmark against
- Don't be surprised if your reviewers start requesting these!

